

SyntacoreTM
Custom cores and tools


Compiler Optimizations Challenges for High-Performance RISC-V Cores

August 25, 2023

Presenter: Sergey Yakushkin

Contributors: Konstantin Vladimirov and Syntacore Compiler Team

info@syntacore.com



SyntacoreTM
Custom cores and tools

Copyright © 2023 Syntacore. All trademarks, product, and brand names belong to their respective owners.



Example High-Performance Core - SCR9

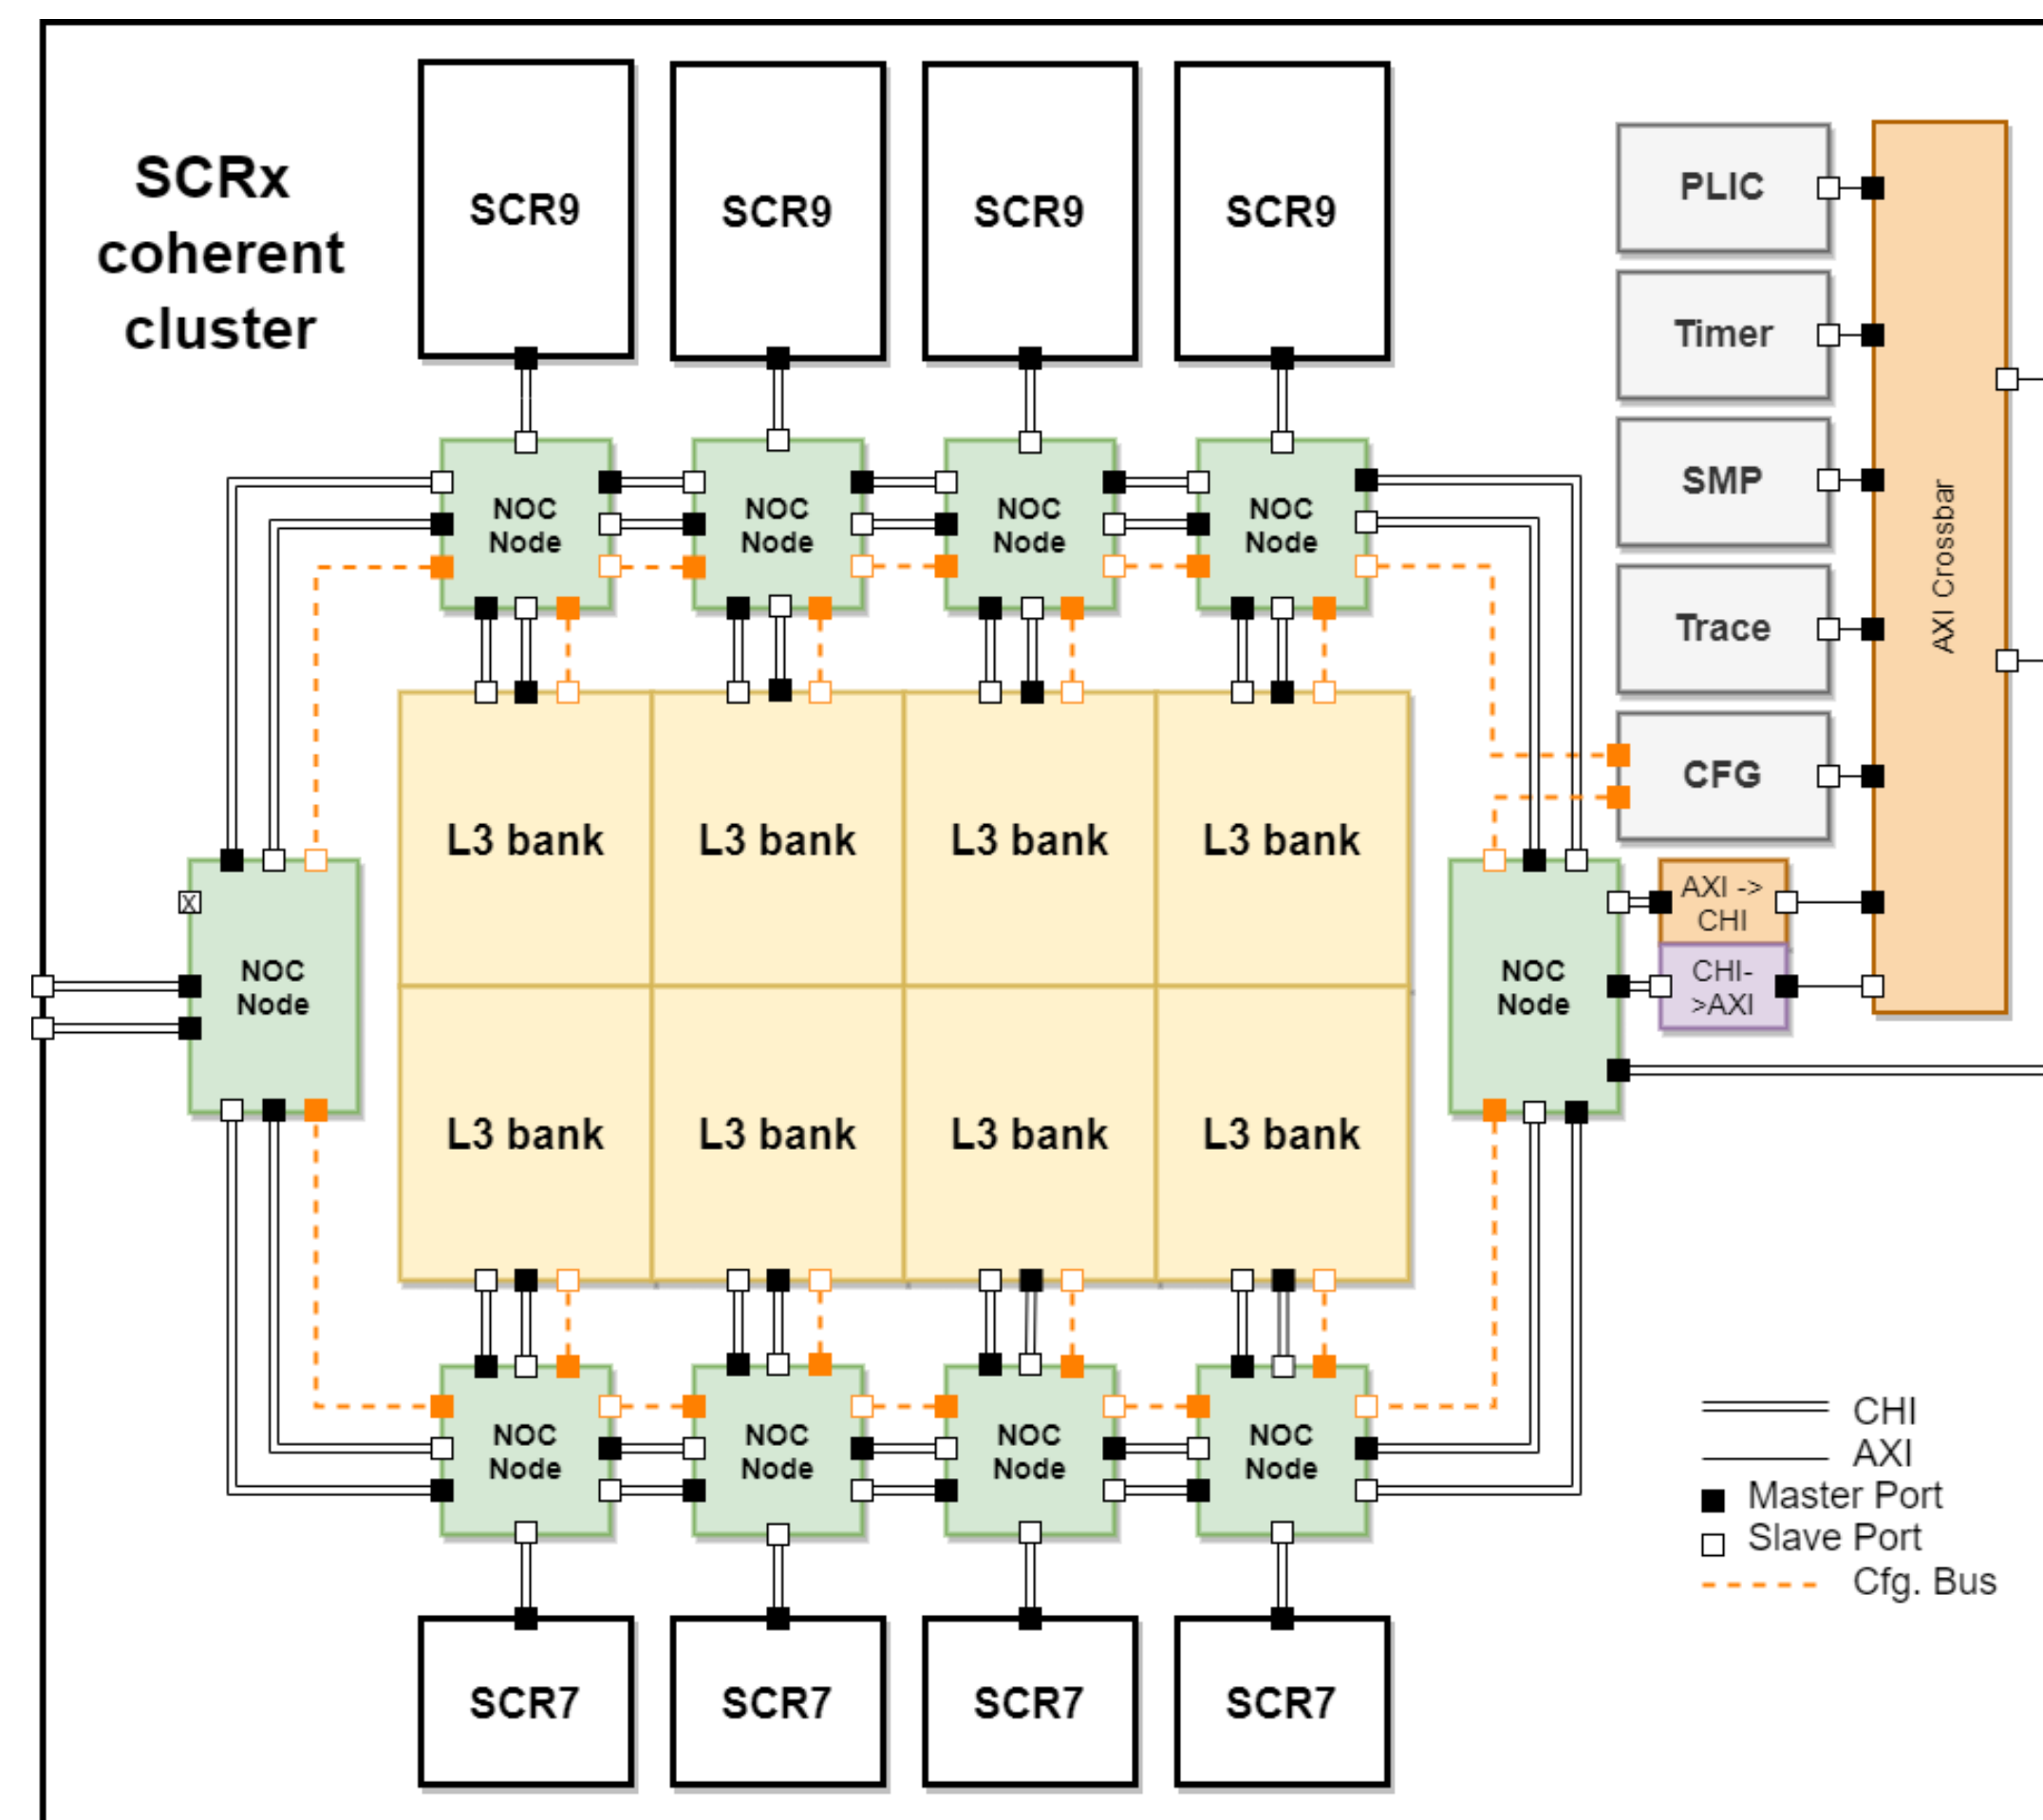
2

Linux-capable application CPU with entry-level server class features:

- 8-16 cores per cluster (SMP and heterogeneous)
- Multi-issue OOO uArch
- Coherent NoC-based L3
- CHI external i/f
- SV39, SV48, SV57
- RVV
- Hypervisor
- AIA
- Accelerators support

Early access program*

(*) some features may be not available in the initial release

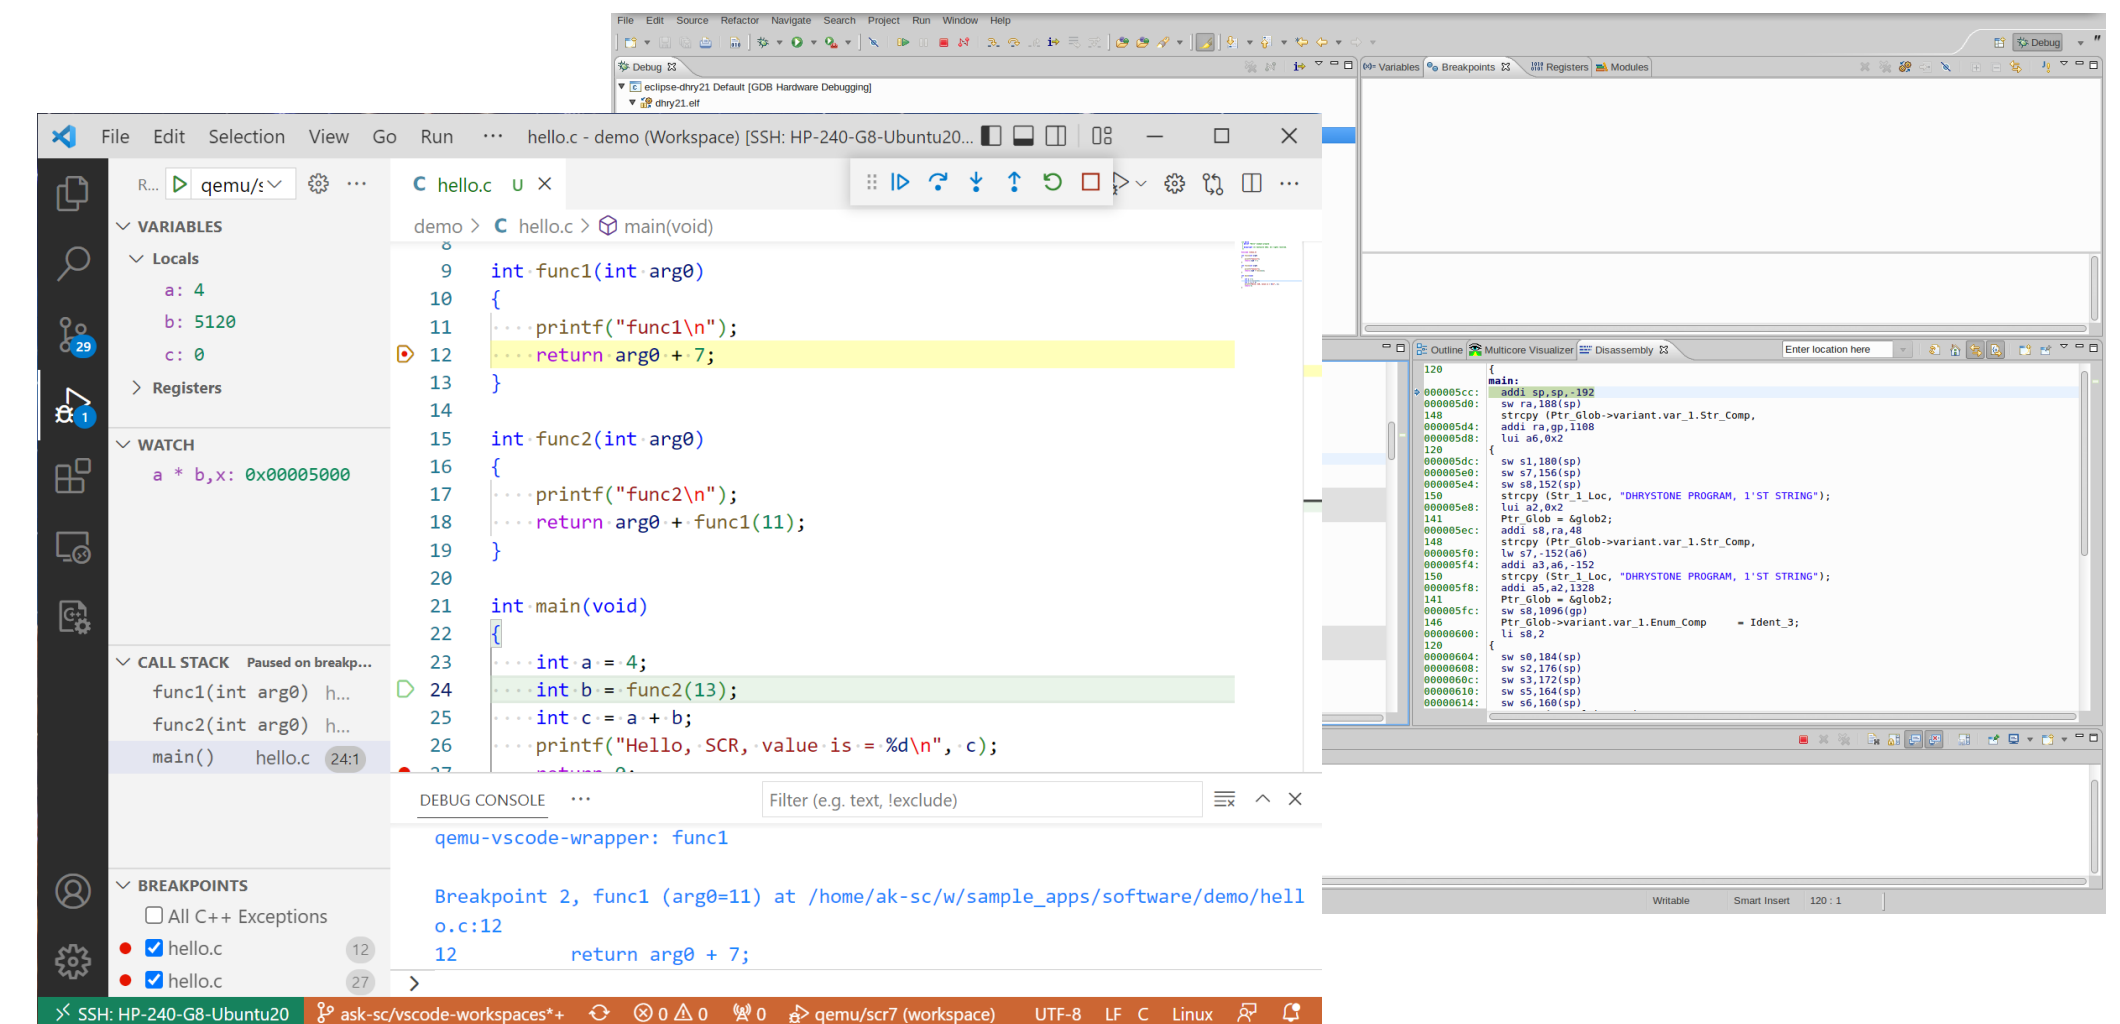



SCR7/9 FPGA-based DevKit

3

Fully-integrated system and DevKit based on

- VCU118 <https://www.xilinx.com/products/boards-and-kits/vcu118.html> or
- UltraScale+ VU19P http://www.hitechglobal.com/Boards/VirtexUltraScale+_VU19P_Board.htm
- Multi-core, up to 24GB RAM, up to 100-150 MHz, 1GB Ethernet, PCI/SSD storage
- Boots upstream Debian Linux kernel 5.15/6.1 LTS
- Integrated toolchain with IDE (supports Bare Metal and Linux targets)
- Extra SW including OpenJDK stable builds



Syntacore Development Toolkit

4

SC-DT 2023.08 release:

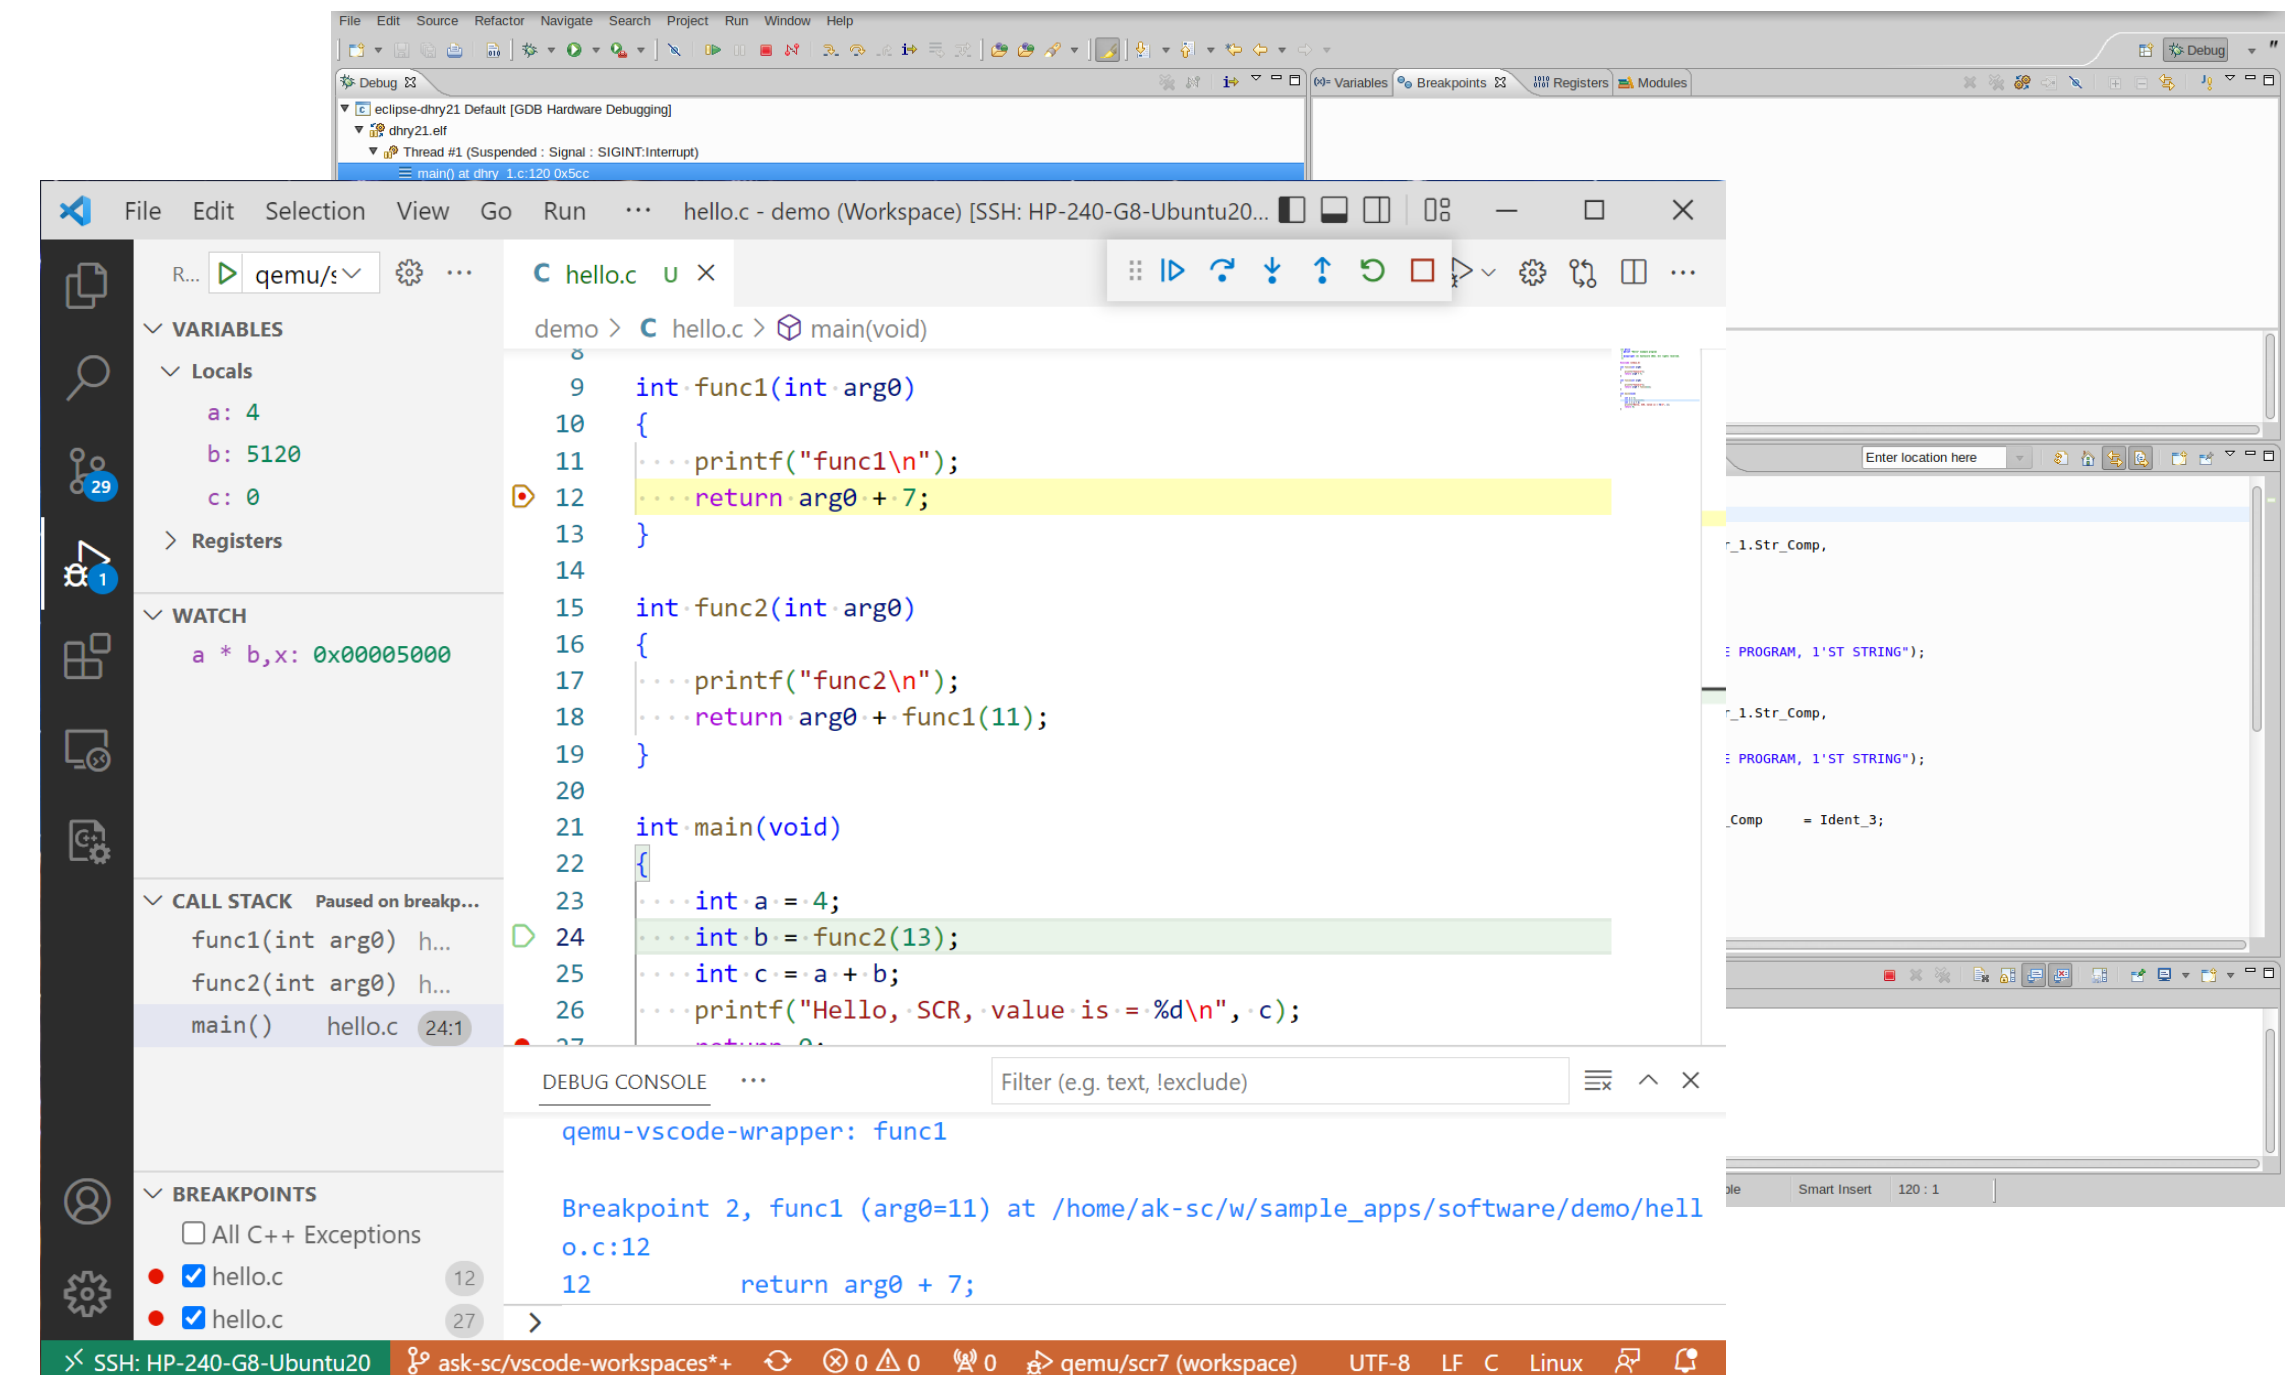
- LLVM 16.x with optimizations
- GNU GDB 13.2
- Open On-Chip Debugger 0.12.x
- QEMU 8.0.3
- GCC 12.2.1
- GNU Binutils 2.38
- Newlib 4.10
- Visual Studio Code and Eclipse



Hosts



Bare Metal + Linux



Also available:

- COMPCERT
- Profiling tools
- OpenJDK

Simulators:

- QEMU
- Spike
- SAIL
- 3rd party vendors

JTAG-based debug solutions:

- Segger J-link
- Olimex ARM-USB-OCD family
- Digilent JTAG-HS2
- more vendors soon



SyntacoreTM
Custom cores and tools

<https://syntacore.com/page/products/sw-tools>

Copyright © 2023 Syntacore. All trademarks, product, and brand names belong to their respective owners.



Upstream LLVM for RISC-V in 2022

- Gaps vs AArch64/GCC, >10% diffs on some workloads
- ABI and conventions are evolving
- Functional issues, e.g. vectorization

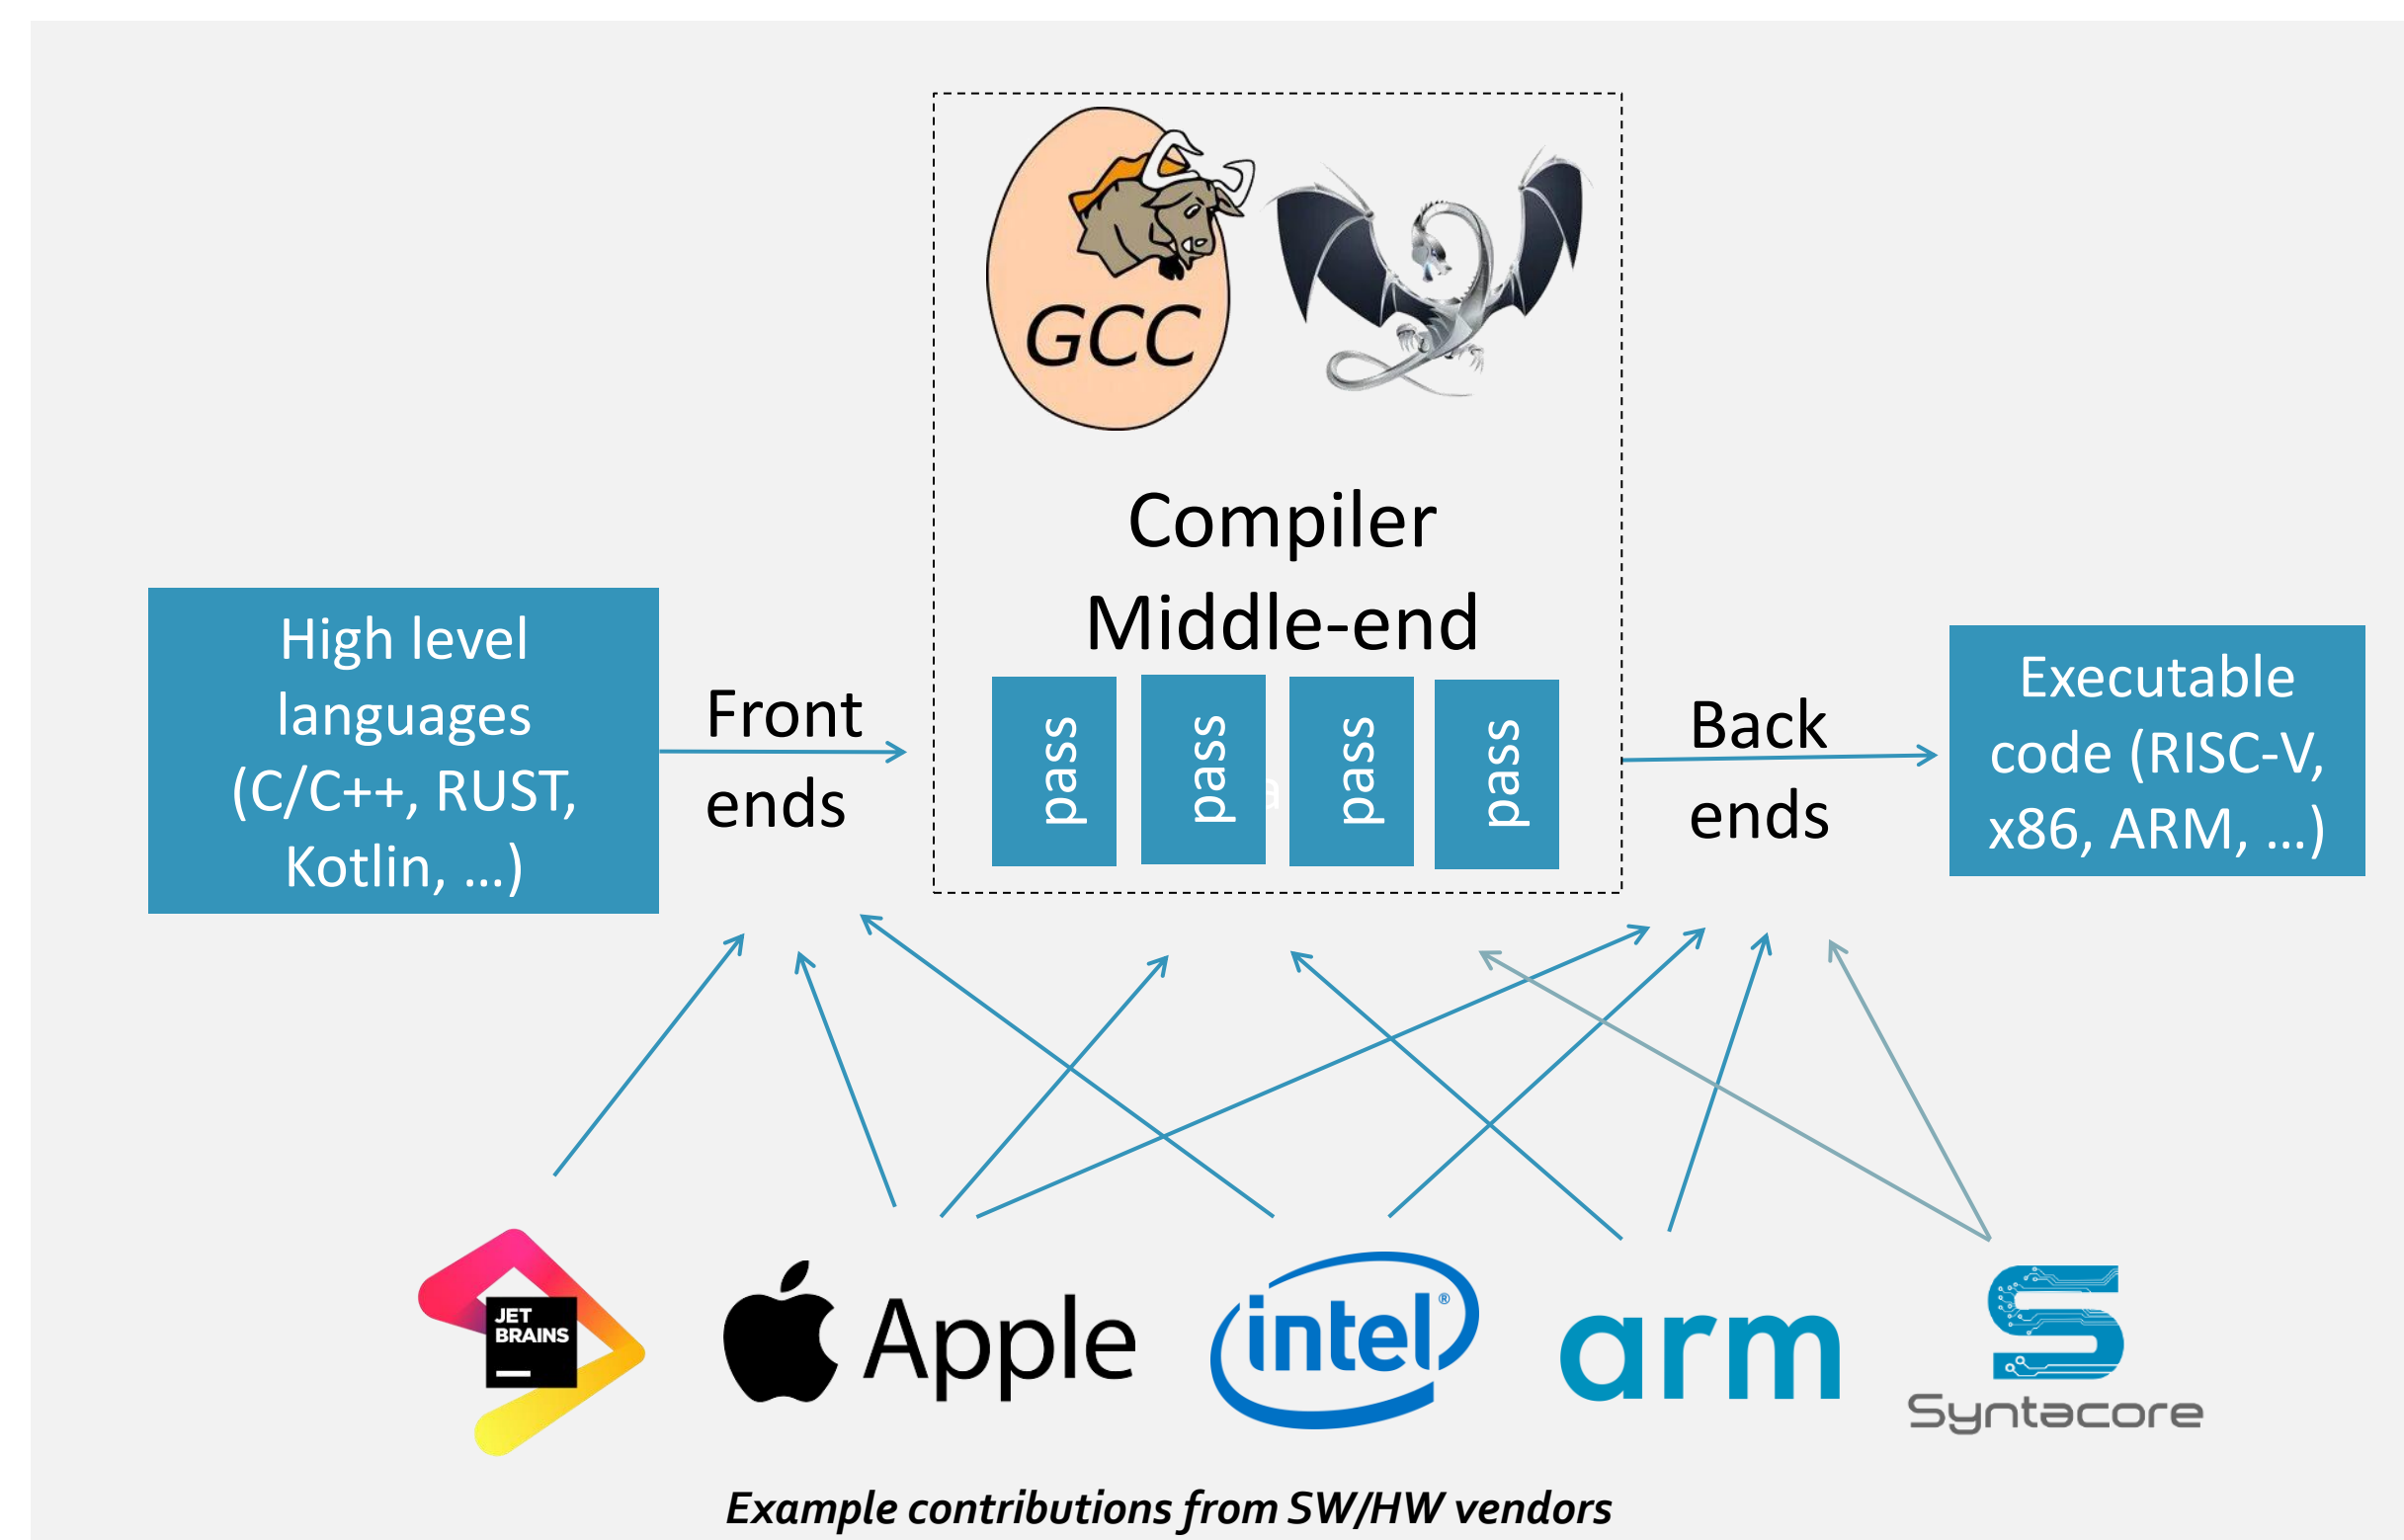
Syntacore LLVM Toolchain

- Improved stability
- SCR uArch-aware optimizations
- Advanced optimizations for RISC-V
- Enhanced RVV auto-vectorization

LLVM is open-source compiler framework:

20+ years, 500 000+ commits, 1500+ developers

primary for OS (Android), languages (Swift, Kotlin, RUST), commercial tools (Intel C++ Compiler), and code analyzers (Coverity)



Loop Optimizations: Unsigned Wrap


Compiler optimization improves codegen for 32-bit loop counters:

- Identify code fragments
- Add dynamic range checks
- Generate multiple specialized code versions for different ranges
- Apply further transformations for optimal code generation

Improves some industry benchmarks up to **18%** (config without Bitmanip)

Patch submitted to LLVM


<https://reviews.llvm.org/D132208>



Source


```
for (unsigned i = 0; i < N; ++i)
  for (unsigned j = 0; j < N; ++j)
    C[i*N+j] = foo();
```

Extra 32/64 conversions inside loop





Dynamic range checks (new)

```
for (unsigned i = 0; i < N; ++i)
  for (unsigned j = 0; j < N; ++j)
    if (overflow(N*N))
      *(C + zext(i*N+j)) = foo();
    else
      /* nuw */ *(C + zext(i*N+j)) = foo();
```




Unswitching

```
if (overflow(N*N))
  two loops with unsigned wrap
else
  for (unsigned i = 0; i < N; ++i)
    for (unsigned j = 0; j < N; ++j)
      *(C + zext(i*N+j /*nuw*/))=foo();
```



IndVar simplify

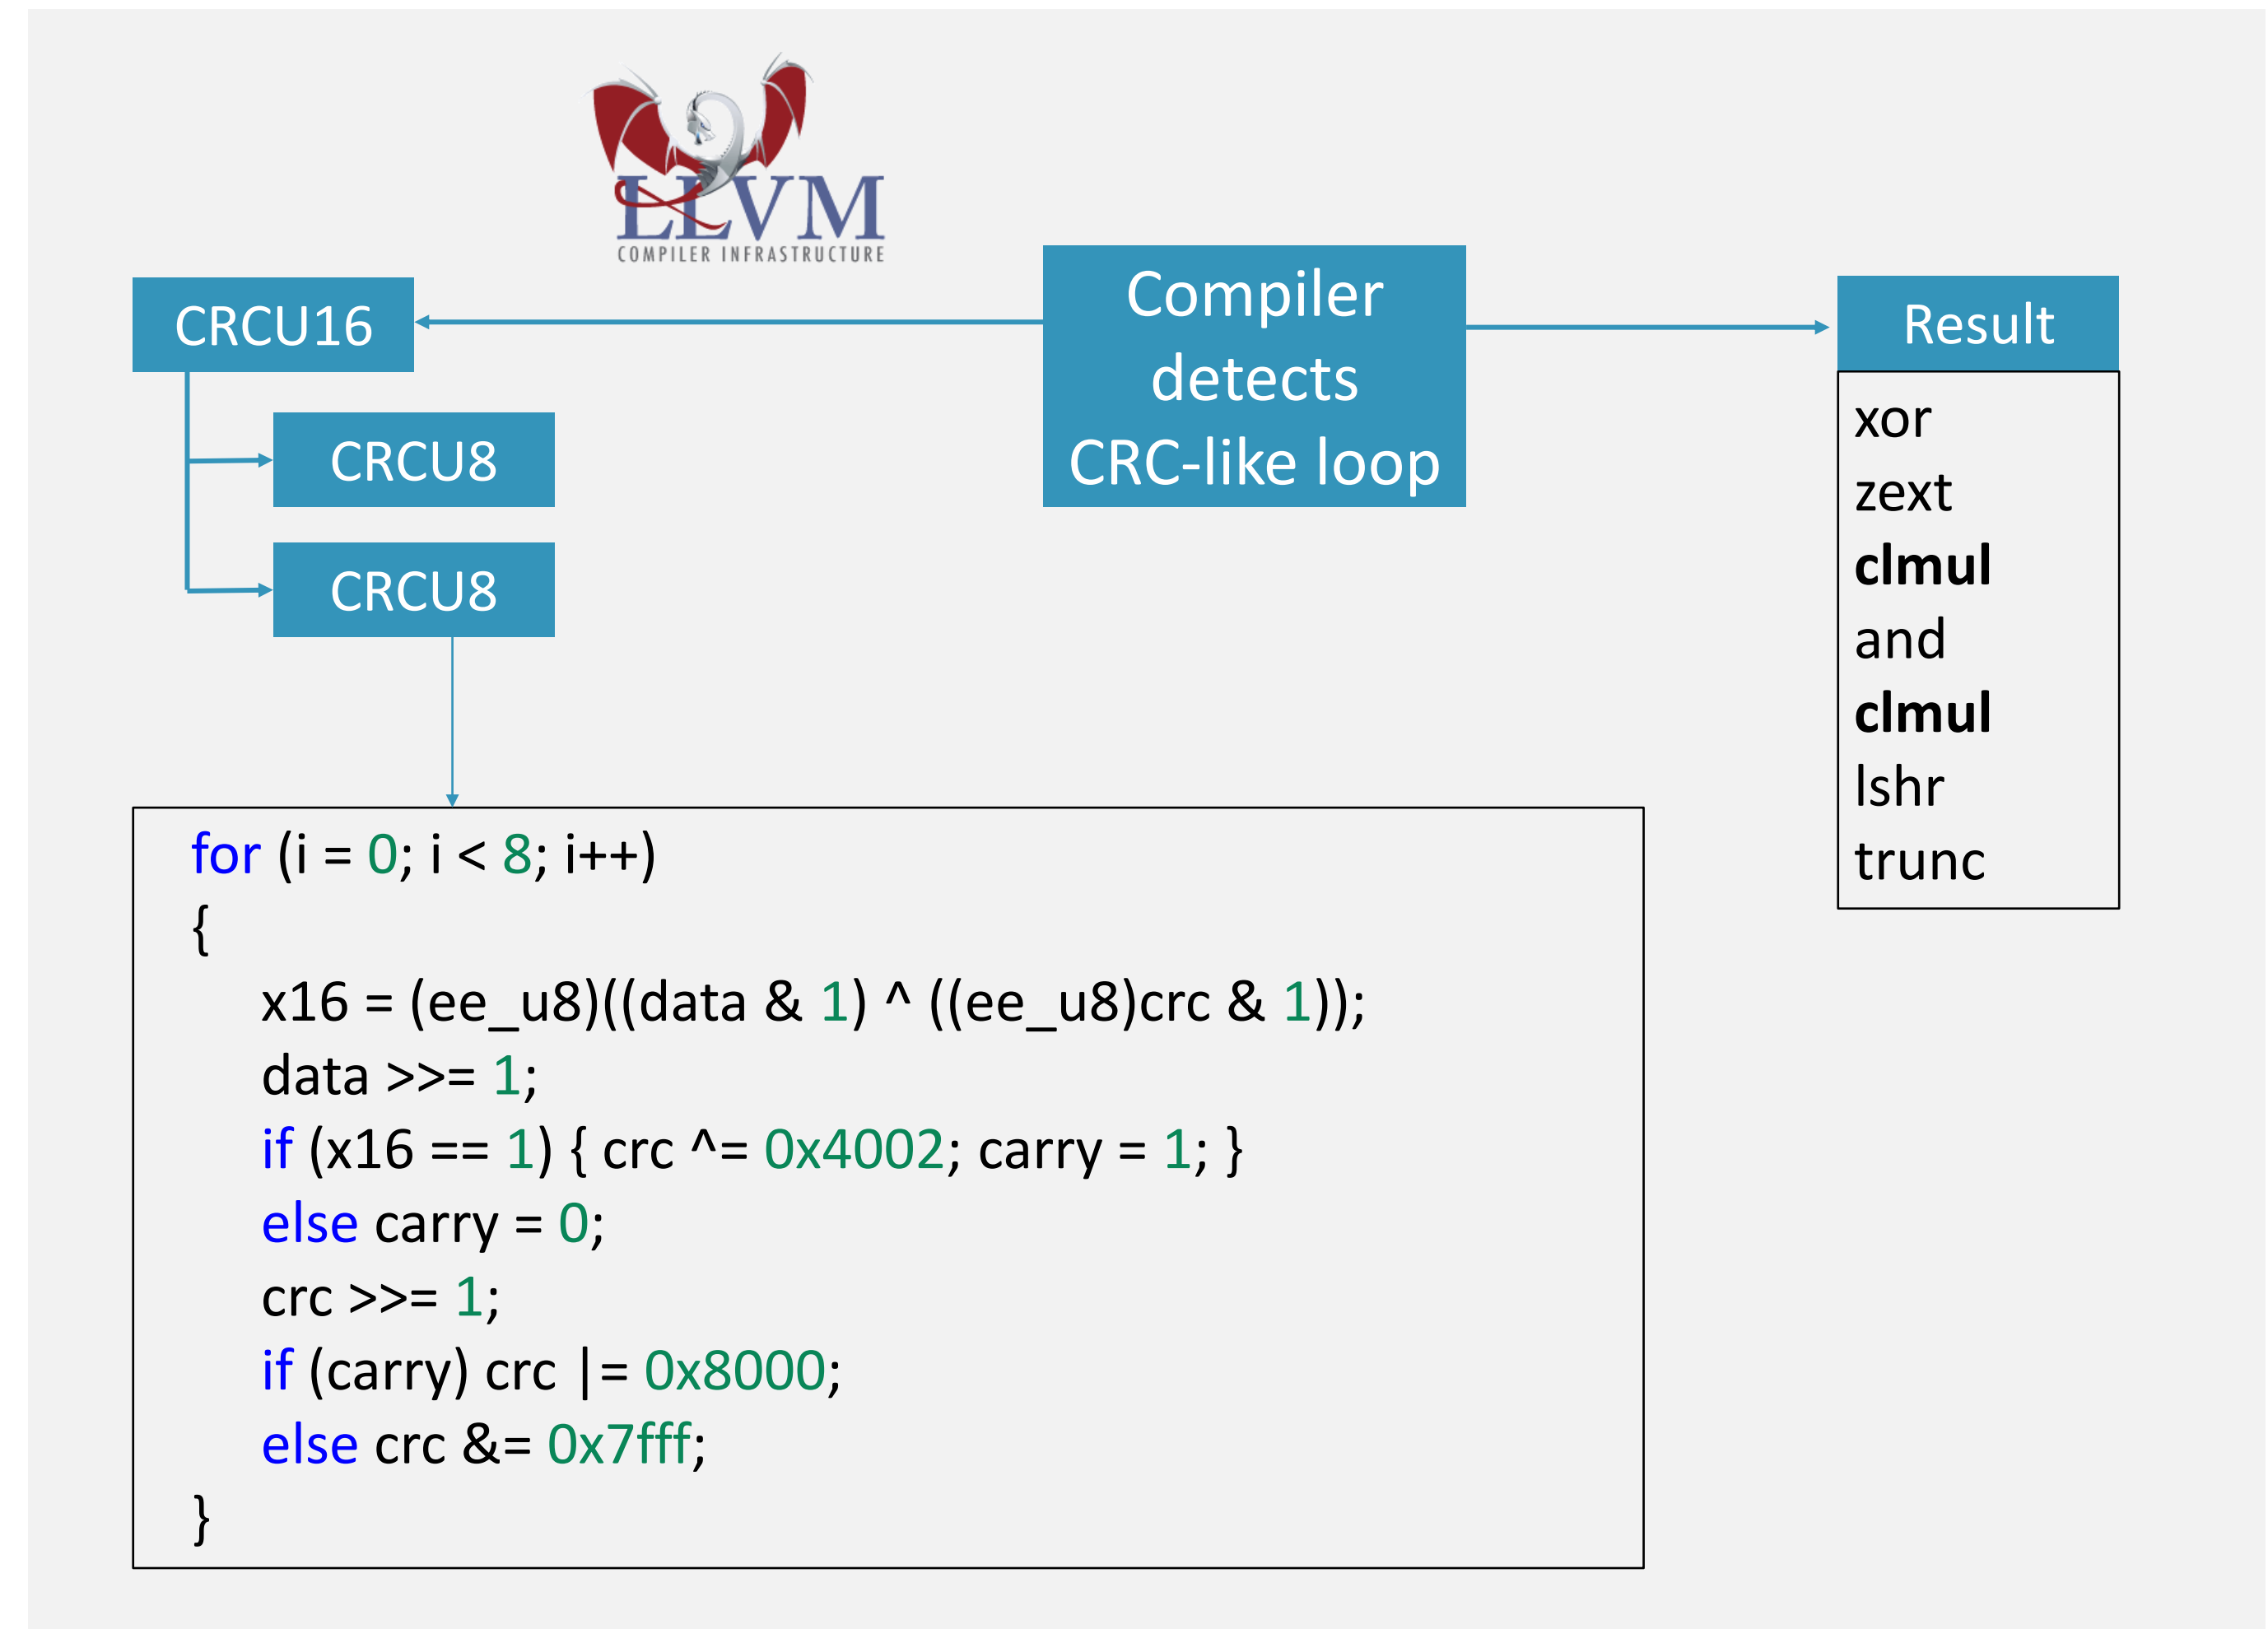
```
if (overflow(N*N))
  two loops with unsigned wrap
else
  for (uint64_t i = 0; i < N; ++i)
    for (uint64_t j = 0; j < N; ++j)
      *(C + (i*N+j)) = foo();
```



Loop Optimizations: CRC-like Patterns

7

- **Compiler** detects loop patterns with generalized CRC computations
- Transforms code and utilize Bitmanip instruction **CLMUL**
- Improves function by 10x, and some industry benchmarks up to **15%**



Relaxation: linker attempts to find shorter instructions

- Example:

```
char arr[42];
char foo(int i) { return arr[i];}
```

- Assembler:

```
lui a5, %hi(arr)      # R_RISCV_HI20, R_RISCV_RELAX
addi a5, a5, %lo(arr) # R_RISCV_L012_I, R_RISCV_RELAX
add  a5, a5, a0
lbu  a0, 0(a5)
ret
```

- Optimization attempt (RISCVMergeBaseOffset):

```
lui a5, %hi(arr)      # R_RISCV_HI20, R_RISCV_RELAX
addi a5, a5, %lo(arr)
add  a5, a5, a0
lbu  a0, %lo(arr)(a5) # R_RISCV_L012_I, R_RISCV_RELAX
ret
```

- Original GP-relaxation **breaks code** in this example

```
lui a5, %hi(arr)      # R_RISCV_HI20, R_RISCV_RELAX
add  a5, a5, a0
lbu  a0, %lo(arr)(a5) # R_RISCV_L012_I, R_RISCV_RELAX
ret
```

- Wrong (original):

```
lui a5, %hi(arr)
add  a5, a5, a0
lbu  a0, offset(gp)
ret
```

- Correct (expected):

```
lui a5, %hi(arr)
add  a5, gp, a0
lbu  a0, offset(a5)
ret
```


- Three new relocation types are supported in Syntacore LLVM toolchain:

- R_RISCV_GPREL_ADD
- R_RISCV_GPREL_LO12_I
- R_RISCV_GPREL_LO12_S

- Generated code example (3 instructions instead of 4):

```
lui a5, %hi(arr)           # R_RISCV_HI20, R_RISCV_RELAX
add a5, a5, a0, %gprel_add(arr) # R_RISCV_GPREL_ADD, R_RISCV_RELAX
lbu a0, %gprel_lo(arr)(a5)  # R_RISCV_GPREL_LO12_I, R_RISCV_RELAX
ret
```

- Similar proposal to fold a non-constant offset into an lui/addi/lw sequence in RISC-V GNU toolchain: [link](#)
- Example impact: ~3% for 445.gobmk SPEC2k

- Partial Redundancy Elimination for Loads: 11 patches ([D141664](#) ... [D143255](#))
 - ~3% on SPEC2k6 471.omnetpp
- Fold terminating condition for any icmp (-lsr-term-fold [D145929](#))
 - ~0.5% overall on SPEC2k6
- RVV Fixes and Improvements
 - Fixed unwinding for RVV spills using DWARF CFA ([D136263](#) [D136264](#))
 - Proposed further optimizations for tail-agnostic policy VMV/VFMV ([D130895](#))

- Syntacore LLVM is close/better than GCC, many patches submitted upstream
 - 10%+ on some SPEC benchmarks
 - Fixed functional bugs, e.g. RVV spill/unwinding
- Areas for further compiler improvements vs AArch64/x86 and GCC
 - Enhancement of generic LLVM passes, MachineCombiner, Peephole, new ISA extensions
 - RVV generic and TA-specific optimizations
 - Loop optimizations and recognition of patterns
- Opportunities for new link-time optimizations with GPREL_* relocations

Thank you!

謝謝



info@syntacore.com



Syntacore[™]
Custom cores and tools

Copyright © 2023 Syntacore. All trademarks, product, and brand names belong to their respective owners.

